

Rohit Verma

+91-90000-00000 | rohit.verma@example.com | linkedin.com/in/rohit-verma-demo |
github.com/rohitverma-demo

Bihar, India
22 September 2026

Hiring Team
StackBlitz

Re: Senior Applied AI Engineer

Dear StackBlitz hiring team,

You're building the core coding agents that power Bolt.new — agent loops, context compression, model routing — and that's exactly the kind of problem I've been solving. Over the last two years I've led the AI systems for a multi-tenant enterprise platform: I built a **LangGraph**-based agentic workflow with human-in-the-loop review that cut research effort by 30%, and I designed a token-metering and LLM observability platform from scratch that cut token costs roughly 90% through prompt caching. I own these systems end-to-end, from architecture to production.

The posting asks for someone who can integrate and evaluate multiple LLM providers, design safe tool-use, and build guardrails. I've worked across Gemini, Mistral, and Groq in production, routing different tasks to different models based on cost and latency trade-offs. My Text-to-SQL engine uses schema-grounded RAG with guardrails against destructive queries and cross-tenant leakage, hitting 75% accuracy across 200+ schemas. I also built a vision-LLM extraction pipeline on Gemini with 90% field-level accuracy. I'm comfortable debugging complex agent behaviors and translating failure patterns into concrete improvements.

StackBlitz is the kind of company I want to be at: a small team shipping to over a million developers, where the patterns for AI-powered development are still being invented. I've mentored five engineers on AI architecture decisions and I'm used to working autonomously in a fast-moving environment. I'd like to bring that experience to Bolt.new's agent team and help push what's possible with in-browser AI development.

Sincerely,
Rohit Verma